



# Probabilistic Changepoint Modeling

An example report for Pittsburgh userR Group

Mikhail Popov

2 December 2018

**Summary** Our imaginary store receives some number of customers per day. We start an advertising campaign which takes some time to have a full effect on the rate of customers we receive daily, and then has a long-term effect after we have stopped it. Using probabilistic programming languages (PPLs), we can specify a Bayesian model and infer the hidden rates.

## Data simulation


Suppose we are operating a store which receives a random number of customers per day, and specifically it's random according to the Poisson distribution with rate  $\lambda_1$ . Then we start an advertisement campaign – which takes a few days to come to full effect – that changes the rate to  $\lambda_2$ . When we stop the ad, it gradually loses effect but the campaign has had a long-lasting effect (e.g. we've gained new customers who come back regularly), which means our store now receives customers at rate  $\lambda_3$ .

Table 1: Simulation parameters

parameter	value
<b>Rates (unknown)</b>	
$\lambda_1$ (old normal)	300
$\lambda_2$ (temporary)	400
$\lambda_3$ (new normal)	350
<b>Other parameters</b>	
$N$ (total days)	180
$T_1$ (ad start)	60
$T_2$ (ad stop)	90
$d_1$ (time for full effect)	7
$d_2$ (time to new normal)	21

In the simulation, we have two weights –  $w_1$  and  $w_2$  – which sum to 1 and produce gradual (albeit not smooth) transitions between the different rates. The change is a slope in this toy scenario, but it can also be a more interesting transition like a sigmoid function such as the [Gompertz curve](#).

Daily customers before, during, and after an advertising campaign



## Modelling

We'll take a look at three similar models –  $\mathcal{M}_1, \mathcal{M}_2, \mathcal{M}_3$  – which model the daily counts of customers  $y_t$  as a [Poisson](#)-distributed random variable with time-varying rate  $\lambda(t)$ , up to a maximum of  $t = N$  days of data:  $y_t \sim \text{Poisson}(\lambda(t))$ ,  $t = 1, \dots, N$ .

The ad campaign starts on  $T_1$  and stops on  $T_2$ . We are interested in inferring  $\lambda_1$  (the rate before  $T_1$ ),  $\lambda_2$  (the rate between  $T_1$  and  $T_2$ ), and  $\lambda_3$  (the rate after  $T_2$ ). We specify  $\lambda \sim \mathcal{N}(300, 100)$  and let Stan implicitly assign a default prior to  $\beta_1$  and  $\beta_2$ .

### Model 1

In the simpler model  $\mathcal{M}_1$ , we ignore the obvious transition periods and model the switch between the rates as immediate:

$$\lambda(t) = \begin{cases} \lambda_1, & \text{if } t \leq T_1, \\ \lambda_2, & \text{if } T_1 < t \leq T_2, \\ \lambda_3, & \text{if } t > T_2. \end{cases}$$

### Model 2

In the slightly more complex model  $\mathcal{M}_2$ , we include the two gradual changes as slopes over  $d_1$  and  $d_2$  days after  $T_1$  and  $T_2$ , respectively. We formalize this as follows:

$$\lambda(t) = \begin{cases} \lambda_1, & \text{if } t \leq T_1, \\ \lambda_2, & \text{if } T_1 + d_1 \leq t \leq T_2, \\ \lambda_3, & \text{if } t \geq T_2 + d_2, \\ \lambda_1 + \beta_1(t - T_1), & \text{if } T_1 < t \leq T_1 + d_1, \\ \lambda_2 + \beta_2(t - T_2), & \text{if } T_2 < t \leq T_2 + d_2. \end{cases}$$

**Note:** in this toy scenario we know  $d_1$  and  $d_2$ , but we can also make them hidden parameters to infer from the data. We could formalize our intuition about their values by assigning them priors  $\mathcal{N}(7, 2)$  and  $\mathcal{N}(14, 4)$ , respectively.

## Model comparison

**Table 2:** Posterior probabilities of the three models given data and the estimated rates, calculated using the `bridgesampling` package.

Model	$\Pr(\mathcal{M} \mathcal{D})$	Estimate (95% Credible Interval)		
		$\lambda_1$ (300)	$\lambda_2$ (400)	$\lambda_3$ (350)
$\mathcal{M}_1$	0.0000	302.9 (298.6–307.1)	387.5 (380.6–394.4)	352.8 (348.9–356.9)
$\mathcal{M}_2$	1.0000	302.8 (298.5–307.1)	395.0 (387.8–402.7)	348.1 (343.8–352.5)

The [Bayes factor](#) (BF) can be used to decide between these two competing models –  $\mathcal{M}_1$  (instant changes between rates) and  $\mathcal{M}_2$  (gradual changes between rates) – by quantifying how much more likely the data  $\mathcal{D}$  is under  $\mathcal{M}_1$  vs  $\mathcal{M}_2$ :

$$\text{BF}_{12} = \frac{p(\mathcal{D}|\mathcal{M}_1)}{p(\mathcal{D}|\mathcal{M}_2)},$$

Using `bridgesampling` (Gronau & Singmann, 2018) allows us to calculate marginal likelihoods –  $p(\mathcal{D}|\mathcal{M})$  – of Stan models really easily and therefore compute the BF (see `?bf`), which – for  $\mathcal{M}_1$  compared to  $\mathcal{M}_2$ , for example – comes out to be 0, meaning there is no evidence for choosing model 1 over model 2.

## Inference results

**Table 3:** Inference using  $\mathcal{M}_2$  (gradual changes between rates).

Parameter	Truth	Point Estimate	95% Credible Interval
<b>Rates</b>			
$\lambda_1$	300	302.8	(298.4, 306.9)
$\lambda_2$	400	394.9	(388.1, 403.0)
$\lambda_3$	350	348.0	(343.9, 352.5)
<b>Differences</b>			
$\lambda_2 - \lambda_1$	100	92.1	(84.1, 100.9)
$\lambda_3 - \lambda_2$	-50	-46.9	(-55.9, -39.1)
$\lambda_3 - \lambda_1$	50	45.2	(38.9, 51.1)

Table 3 shows the inferred differences of rates and we can see that our imaginary advertisement had a positive, statistically significant impact on how many imaginary customers our imaginary store receives per day on average, both during the campaign and long after.

## References

- Allaire, J., Xie, Y., McPherson, J., Luraschi, J., Ushey, K., Atkins, A., ... Chang, W. (2018). Rmarkdown: Dynamic documents for r. Retrieved from <https://CRAN.R-project.org/package=rmarkdown>
- Bache, S. M. (2015). Import: An import mechanism for r. Retrieved from <https://CRAN.R-project.org/package=import>
- Gronau, Q. F., & Singmann, H. (2018). Bridgesampling: Bridge sampling for marginal likelihoods and bayes factors. Retrieved from <https://CRAN.R-project.org/package=bridgesampling>
- Henry, L., & Wickham, H. (2018). Purrr: Functional programming tools. Retrieved from <https://CRAN.R-project.org/package=purrr>
- Müller, K. (2017). Here: A simpler way to find your files. Retrieved from <https://CRAN.R-project.org/package=here>
- Popov, M. (2018). Wmfpar: Wikimedia foundation's product analytics reporting template. Retrieved from <https://github.com/bearloga/wmf-product-analytics-report>
- R Core Team. (2018). R: A language and environment for statistical computing. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from <https://www.R-project.org/>
- Stan Development Team. (2018). RStan: The R interface to Stan. Retrieved from <http://mc-stan.org/>
- Wickham, H. (2016). Ggplot2: Elegant graphics for data analysis. Springer-Verlag New York. Retrieved from <http://ggplot2.org>
- Wickham, H., François, R., Henry, L., & Müller, K. (2018). Dplyr: A grammar of data manipulation. Retrieved from <https://CRAN.R-project.org/package=dplyr>
- Wickham, H., & Henry, L. (2018). Tidyr: Easily tidy data with 'spread()' and 'gather()' functions. Retrieved from <https://CRAN.R-project.org/package=tidyr>
- Xie, Y. (2014). Knitr: A comprehensive tool for reproducible research in R. In V. Stodden, F. Leisch, & R. D. Peng (Eds.), Implementing reproducible computational research. Chapman; Hall/CRC. Retrieved from <http://www.crcpress.com/product/isbn/9781466561595>
- Xie, Y. (2015). Dynamic documents with R and knitr (2nd ed.). Boca Raton, Florida: Chapman; Hall/CRC. Retrieved from <https://yihui.name/knitr/>
- Xie, Y. (2018). Knitr: A general-purpose package for dynamic report generation in r. Retrieved from <https://yihui.name/knitr/>
- Zhu, H. (2018). KableExtra: Construct complex table with 'kable' and pipe syntax. Retrieved from <https://CRAN.R-project.org/package=kableExtra>
- Zhu, H., Tsai, T., & Trivison, T. (2018). Memor: A 'rmarkdown' template that can be highly customized. Retrieved from <https://CRAN.R-project.org/package=memor>